


Pattern Recognition Systems, Problem Solving using Pattern Recognition

Face Recognition

A Classification Problem



Authors: team **NTY**

Ng Siew Pheng (A0198525R)

Tea Lee Seng (A0198538J)

Yang Xiaoyan (A0056720L)

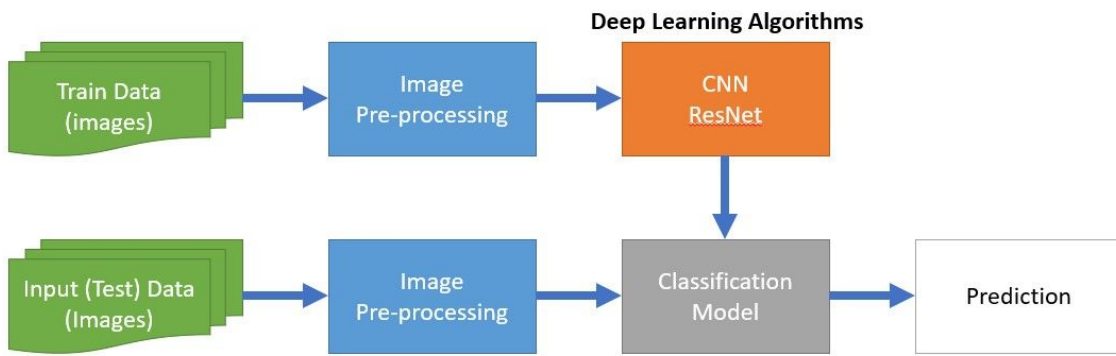
Table of contents	
Introduction	3
Train, Validation and Test Strategy	4
Data preprocessing	5
Deep learning algorithms	6
Image Data Generator	6
Convolutional Neural Network (CNN)	7
Training with full images	8
Training with face only images	8
Residual neural network (ResNet)	11
Training with full images	12
Training with face only images	13
Model Test Result Comparison	17
Challenges	18
Videos to images	18
The usage of jupyter notebook and code quality	18
Slower training due to single process image augmentation	20
ImageDataGenerator.flow_from_dataframe(), implicit shuffle=True	23
High validation accuracy BUT Low testing accuracy	24
Conclusion	24
References:	25

1. Introduction

A facial recognition system is a technology capable of identifying or verifying a person from a digital image or a video frame from a video source. Facial recognition is a Biometric Artificial Intelligence based application that can uniquely identify a person by analysing patterns based on the person's facial textures and shapes.

Face recognition can be applied Security, Healthcare, Marketing etc areas.

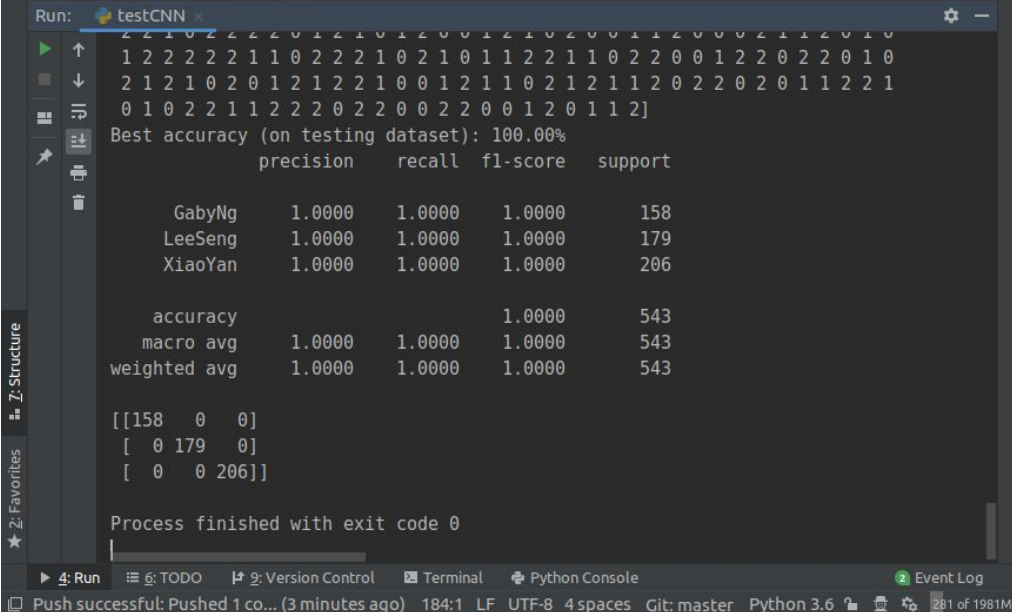
In this project, we will build a facial recognition system by applying deep learning algorithms to identify three people's faces. We will be looking at using CNN and ResNet and determine which model is better.



2. Train, Validation and Test Strategy

The data set we are using for this image classification problem, are 2D face images from 3 different people. We record selfie videos of each person from different angles. Each video file is tagged to a particular person. There are about 1800 to 2500 images for each person to perform training. The videos can be downloaded from <http://home.leeseng.tech/rawData-20190910.zip>

Initially, we apply 70% training, 20% validation and 10% testing data strategy. The result is too good to be true for the 10% testing data.



```
Run: testCNN
2 2 1 0 2 2 2 0 1 2 1 0 1 2 0 0 1 2 1 0 2 0 0 1 1 2 0 0 0 2 1 1 2 0 1 0
1 2 2 2 2 1 1 0 2 2 2 1 0 2 1 0 1 1 2 2 1 1 0 2 2 0 0 1 2 2 0 2 2 0 1 0
2 1 2 1 0 2 0 1 2 1 2 2 1 0 0 1 2 1 1 0 2 1 2 1 1 2 0 2 2 0 2 0 1 1 2 2 1
0 1 0 2 2 1 1 2 2 2 0 2 2 0 0 2 2 0 0 1 2 0 1 1 2]
Best accuracy (on testing dataset): 100.00%
      precision    recall  f1-score   support

   GabyNg         1.0000         1.0000         1.0000         158
   LeeSeng         1.0000         1.0000         1.0000         179
   XiaoYan         1.0000         1.0000         1.0000         206

 accuracy          1.0000          1.0000          1.0000         543
 macro avg          1.0000          1.0000          1.0000         543
weighted avg          1.0000          1.0000          1.0000         543

[[158  0  0]
 [  0 179  0]
 [  0  0 206]]

Process finished with exit code 0
```

We suspect the 10% test data is too similar due to the case that we are using data from videos.

A separate set of 21 images, from the same 3 people, were used for testing. And these test images are not part of the same video image series. For example, the background might be totally different. This will allow us to test the accuracy of our model when encountering data that it have not been trained with.

3. Data preprocessing

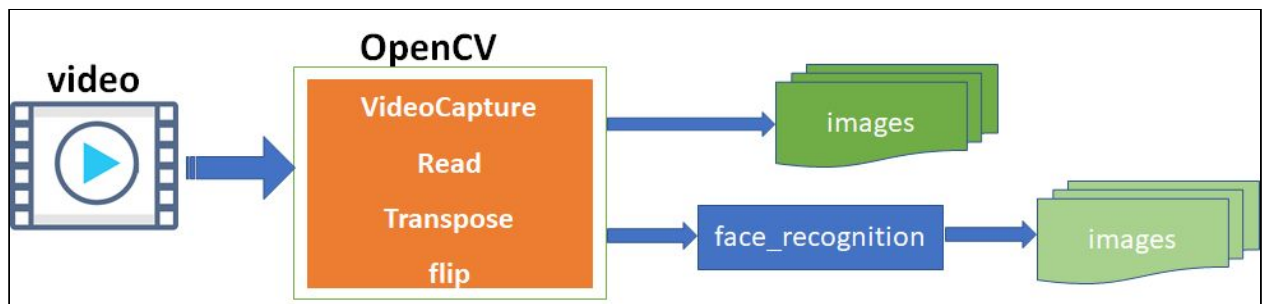
Before we can start training our model, we need to perform data pre-processing to our input data.

We extracted our images from 2 source formats, namely video (MP4) format and photo (JPEG) format. In order to extract 2D images from the video at a specific frame per second, we created a pre-processing utility, based on OpenCV library. The extracted images are captured in folder, data-full.

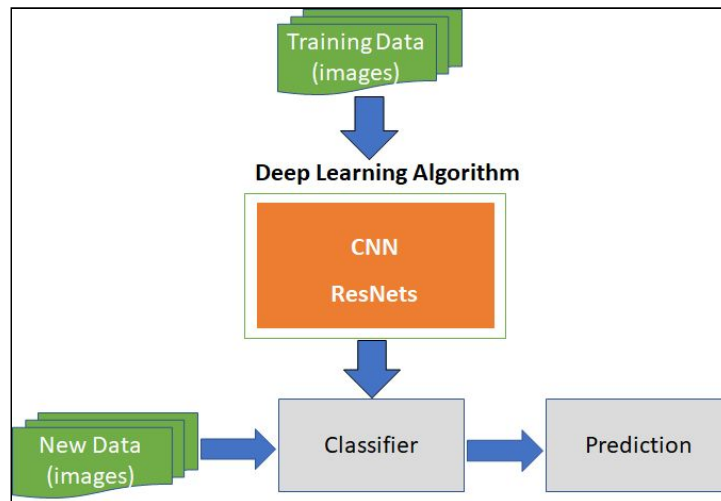
In addition to using these 2D images, we also utilized library, [face_recognition](#), to crop the face from these 2D images. But there is a small problem. We noticed that if we tried detecting the face using the 2D image in its existing orientation, we weren't able to detect and extract the face. This seems to be due to mobile phone model that we used to capture our own selfie videos. Face clipping data was decided due to less than satisfactory result from normal selfie videos.

Thus, in the face extraction utility, [videoToImages.py](#), if we are unable to detect any faces at 0 degrees, we will rotate the image by 90 degrees and try to detect the face detection. This step is repeated until either a face is detected or the image has completed a full rotation. The extracted face images are captured into folder, data-face. Full training/testing set, with minor cleaning from wrong face cropping, can be downloaded at

<http://home.leeseng.tech/training-data-20190925.zip> 2.7GB



4. Deep learning algorithms



4.1. Image Data Generator

Before loading the image data to the deep learning algorithms, we need to perform data augmentation on the training images.

We will perform data augmentation using 3 different Image Data Generators, as shown below, to determine which will help train our model, to yield the highest and consistent accuracy.

Parameter	Rescale Generator	Samplewise Generator	Samplewise + Rescale Generator
samplewise_center		True	True
samplewise_std_normalization		True	True
rescale	1.0/255		1.0/255
width_shift_range	0.1	0.1	0.1
height_shift_range	0.1	0.1	0.1
rotation_range	20	20	20
zoom_range	0.1	0.1	0.1
shear_range	0.15	0.15	0.15
horizontal_flip	True	True	True
vertical_flip	False	False	False
fill_mode	nearest	nearest	nearest

4.2. Convolutional Neural Network (CNN)

We have constructed the following CNN model and using an input image of pixel 128 x 128.

Model: "sequential"		
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 126, 126, 32)	896
activation (Activation)	(None, 126, 126, 32)	0
conv2d_1 (Conv2D)	(None, 124, 124, 64)	18496
activation_1 (Activation)	(None, 124, 124, 64)	0
max_pooling2d (MaxPooling2D)	(None, 62, 62, 64)	0
dropout (Dropout)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73856
activation_2 (Activation)	(None, 60, 60, 128)	0
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 128)	0
conv2d_3 (Conv2D)	(None, 28, 28, 256)	295168
activation_3 (Activation)	(None, 28, 28, 256)	0
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 256)	0
dropout_1 (Dropout)	(None, 14, 14, 256)	0
conv2d_4 (Conv2D)	(None, 12, 12, 256)	590080
activation_4 (Activation)	(None, 12, 12, 256)	0
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 256)	0
dropout_2 (Dropout)	(None, 6, 6, 256)	0
conv2d_5 (Conv2D)	(None, 4, 4, 512)	1180160
activation_5 (Activation)	(None, 4, 4, 512)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 512)	4194816
activation_6 (Activation)	(None, 512)	0
dropout_3 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 3)	1539
Total params: 6,355,011		
Trainable params: 6,355,011		
Non-trainable params: 0		

With this model, we will start training it using the 3 different image data generators, as mentioned in section 5.1.

4.2.1. Training with full images

- Rescale ImageDataGenerator

```
testout: [0 0 0 0 2 2 2 2 2 2 2 2 1 1 1 1 1 1]
predout: [1 0 1 0 0 1 2 2 1 1 2 2 1 0 1 1 2 2 1 2 0]
Best accuracy (on testing dataset): 47.62%
      precision    recall  f1-score   support

   GabyNg      0.6000      0.6000      0.6000         5
   LeeSeng      0.3333      0.4286      0.3750         7
   XiaoYan      0.5714      0.4444      0.5000         9


 accuracy      0.4762
macro avg      0.5016      0.4910      0.4917         21
weighted avg      0.4989      0.4762      0.4821         21

[[3 2 0]
 [1 3 3]
 [1 4 4]]
```

4.2.2. Training with face only images

- Rescale ImageDataGenerator

Training Model : [CNN FaceTrain Rescale.ipynb](#)




```

test result for CNN using face only images in rescale data generator
Batch size is: 128
testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1]
predout: [1 0 0 0 0 2 2 2 2 2 2 2 2 2 0 1 0 1 1 2 1]
Best accuracy (on test data set): 80.95%

```

	precision	recall	f1-score	support
GabyNg	0.6667	0.8000	0.7273	5
LeeSeng	0.8333	0.7143	0.7692	7
XiaoYan	0.8889	0.8889	0.8889	9
accuracy			0.8095	21
macro avg	0.7963	0.8011	0.7951	21
weighted avg	0.8175	0.8095	0.8105	21

```

[[4 1 0]
 [1 5 1]
 [1 0 8]]

```

Based on the result, with batch size set as 128, we wonder what will the impact be we use a lower batch size. The result below shows that accuracy has dropped.

```

testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1]
predout: [0 0 0 2 1 2 0 0 2 1 2 2 0 0 1 0 1 1 1 2 1]
Batch size is 64
Best accuracy (on testing dataset): 57.14%

```

	precision	recall	f1-score	support
GabyNg	0.3750	0.6000	0.4615	5
LeeSeng	0.7143	0.7143	0.7143	7
XiaoYan	0.6667	0.4444	0.5333	9
accuracy			0.5714	21
macro avg	0.5853	0.5862	0.5697	21
weighted avg	0.6131	0.5714	0.5766	21


```

[[3 1 1]
 [1 5 1]
 [4 1 4]]

```

With this result, we have decided to set batch_size for our image data generator to 128.

- Samplewise ImageDataGenerator
Training Model : [CNN FaceTrain Sample.ipynb](#)




```
Test Results for CNN using face only images in samplewise data generator
testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1]
predout: [0 0 0 0 0 2 2 0 2 2 2 2 2 0 1 1 1 1 1 0 1]
Batch size is 128
Best accuracy (on testing dataset): 85.71%
      precision    recall  f1-score   support

   GabyNg      0.6250      1.0000      0.7692         5
   LeeSeng      1.0000      0.8571      0.9231         7
   XiaoYan      1.0000      0.7778      0.8750         9

 accuracy              0.8571         21
 macro avg           0.8750      0.8783      0.8558         21
weighted avg           0.9107      0.8571      0.8658         21

[[5 0 0]
 [1 6 0]
 [2 0 7]]
```

- Samplewise and Rescale ImageDataGenerator
Training Model : [CNN FaceTrain Sample Rescale.ipynb](#)



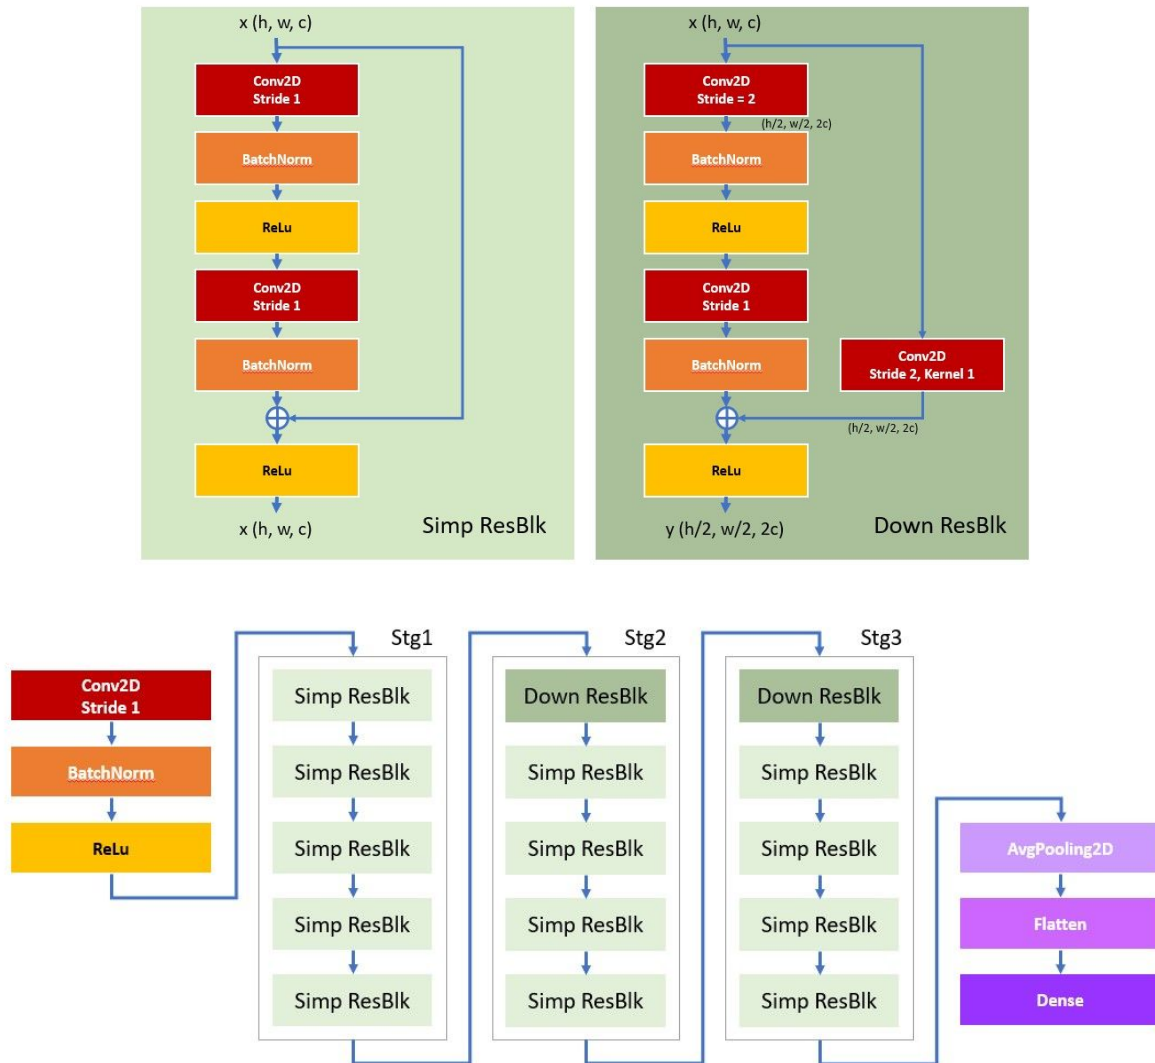
```
Test Results for CNN using face only images in samplewise
and rescale data generator
testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1]
predout: [1 0 0 0 0 2 2 2 2 2 2 2 2 2 0 1 1 0 1 1 0 1]
Batch size is 128
Best accuracy (on testing dataset): 80.95%
```

	precision	recall	f1-score	support
GabyNg	0.5714	0.8000	0.6667	5
LeeSeng	0.8333	0.7143	0.7692	7
XiaoYan	1.0000	0.8889	0.9412	9
accuracy			0.8095	21
macro avg	0.8016	0.8011	0.7924	21
weighted avg	0.8424	0.8095	0.8185	21

```
[[4 1 0]
 [2 5 0]
 [1 0 8]]
```

4.3. Residual neural network (ResNet)

We have constructed the following ResNet model and using an input image of pixel 128 x 128.



Total params: 473,411
Trainable params: 471,139
Non-trainable params: 2,272

4.3.1. Training with full images

- Rescale ImageDataGenerator

Training Model: [ResNet_FullTrain_Rescale.ipynb](#)

```
testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1]
predout: [1 0 1 1 1 2 2 2 2 1 1 1 0 2 1 1 1 1 1 0 1]
Best accuracy (on testing dataset): 57.14%
      precision    recall  f1-score   support


   GabyNg      0.3333      0.2000      0.2500         5
   LeeSeng      0.4615      0.8571      0.6000         7
   XiaoYan      1.0000      0.5556      0.7143         9

   accuracy              0.5714         21
  macro avg      0.5983      0.5376      0.5214         21
 weighted avg      0.6618      0.5714      0.5656         21

[[1 4 0]
 [1 6 0]
 [1 3 5]]
```

- Samplewise ImageDataGenerator

Training Model: [ResNet_FullTrain_Sample.ipynb](#)



```
testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1]
predout: [0 0 1 1 1 2 1 2 1 1 1 1 0 0 1 0 1 0 1 2 1]
Best accuracy (on testing dataset): 38.10%
      precision    recall  f1-score   support

   GabyNg      0.3333      0.4000      0.3636         5
   LeeSeng      0.3333      0.5714      0.4211         7
   XiaoYan      0.6667      0.2222      0.3333         9

   accuracy              0.3810         21
  macro avg      0.4444      0.3979      0.3727         21
 weighted avg      0.4762      0.3810      0.3698         21

[[2 3 0]
 [2 4 1]
 [2 5 2]]
```

- Samplewise and Rescale ImageDataGenerator


```
testout: [0 0 0 0 0 1 2 2 2 2 2 2 2 2]
predout: [1 0 0 1 0 1 2 0 2 2 1 1 1 0]
Best accuracy (on testing dataset): 46.67%
```

	precision	recall	f1-score	support
GabyNg	0.5000	0.6000	0.5455	5
LeeSeng	0.1667	1.0000	0.2857	1
XiaoYan	1.0000	0.3333	0.5000	9
accuracy			0.4667	15
macro avg	0.5556	0.6444	0.4437	15
weighted avg	0.7778	0.4667	0.5009	15

```
[[3 2 0]
 [0 1 0]
 [3 3 3]]
```

4.3.2. Training with face only images

- Rescale ImageDataGenerator
Training Model: [ResNet_FaceTrain_Rescale.ipynb](#)



Using the input test images in its original form to perform testing.

```

testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1]
predout: [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 1]
Batch size is 128
Best accuracy (on testing dataset): 23.81%

```

	precision	recall	f1-score	support
GabyNg	0.0000	0.0000	0.0000	5
LeeSeng	0.2632	0.7143	0.3846	7
XiaoYan	0.0000	0.0000	0.0000	9
accuracy			0.2381	21
macro avg	0.0877	0.2381	0.1282	21
weighted avg	0.0877	0.2381	0.1282	21

```

[[0 5 0]
 [2 5 0]
 [0 9 0]]

```

Tested using resultant images, generated after running face extraction utility on the input test images

```

testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1]
predout: [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1]
Batch size is 128
Best accuracy (on testing dataset): 33.33%

```


	precision	recall	f1-score	support
GabyNg	0.0000	0.0000	0.0000	5
LeeSeng	0.3333	1.0000	0.5000	7
XiaoYan	0.0000	0.0000	0.0000	9
accuracy			0.3333	21
macro avg	0.1111	0.3333	0.1667	21
weighted avg	0.1111	0.3333	0.1667	21

```

[[0 5 0]
 [0 7 0]
 [0 9 0]]

```

- Samplewise ImageDataGenerator
Training Model: [ResNet FaceTrain Sample.ipynb](#)




```
Test Results for ResNet using face only images in samplewise data generator
testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1]
predout: [1 0 0 0 0 2 2 2 2 2 0 0 2 0 1 1 0 1 1 0 1]
Batch size is 128
Best accuracy (on testing dataset): 66.67%
```

	precision	recall	f1-score	support
GabyNg	0.4000	0.8000	0.5333	5
LeeSeng	0.8333	0.7143	0.7692	7
XiaoYan	1.0000	0.5556	0.7143	9
accuracy			0.6667	21
macro avg	0.7444	0.6899	0.6723	21
weighted avg	0.8016	0.6667	0.6895	21

```
[[4 1 0]
 [2 5 0]
 [4 0 5]]
```


- Samplewise and Rescale ImageDataGenerator
Training Model: [ResNet FaceTrain Sample and Rescale.ipynb](#)



```
testout: [0 0 0 0 0 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1]
predout: [2 0 0 0 0 2 0 0 1 0 0 0 2 0 1 2 1 2 1 0 1]
Batch size is 128
Best accuracy (on testing dataset): 47.62%
```

	precision	recall	f1-score	support
GabyNg	0.3636	0.8000	0.5000	5
LeeSeng	0.8000	0.5714	0.6667	7
XiaoYan	0.4000	0.2222	0.2857	9
accuracy			0.4762	21
macro avg	0.5212	0.5312	0.4841	21
weighted avg	0.5247	0.4762	0.4637	21

```
[[4 0 1]
 [1 4 2]
 [6 1 2]]
```

4.4. Model Test Result Comparison

Using a set of 21 full 2D images as test set, we performed prediction using the models we have trained so far. Following are the test results we obtained.

Model		Train Data	Test Data	
			Full Image	Face Only
CNN	rescale	Full Image	47.6%	
		Face Only	47.62%	80.95%
	Samplewise	Full Image		
		Face Only	28.57%	85.7%, 69.57%
	Samplewise+rescale	Full Image		
		Face Only	33.33%	76.19%
ResNet	rescale	Full Image	57%	
		Face Only	23.81%	33.3%
	Samplewise	Full Image	67%	
		Face Only	42.86%	66.67%
	Samplewise+rescale	Full Image	47%	
		Face Only	61.90%	47.62%

Observation:

- Using CNN Model and data augmentation on the images to detect and crop out the face image and applying samplewise setting in the data generator, yielded the best accuracy of 85.7%.
- Whereas using full image for both train and test yielded the worst result. Interesting the worst prediction is coming from classification of child pictures in our data sets.
- For the resNet, we are getting better test results with sample wise setting (samplewise_center=True, samplewise_std_normalization=True,) as compared with using rescale (rescale = 1. / 255) the images.

5. Challenges

5.1. Videos to images

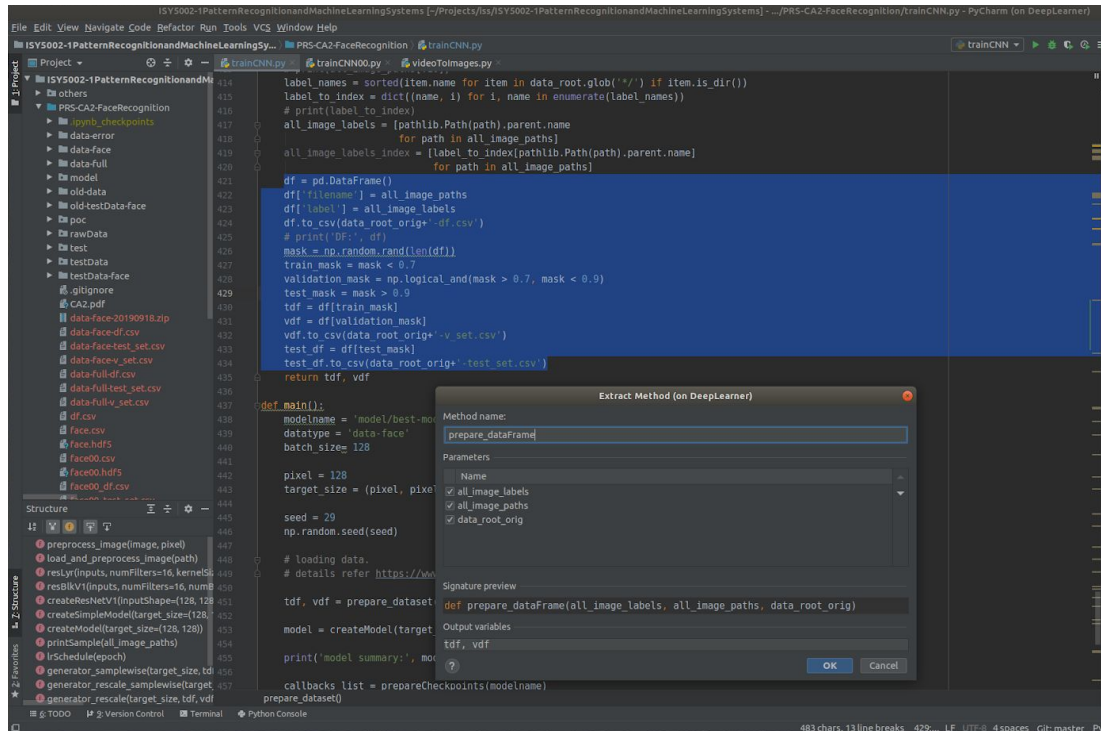
We used mobile phones to capture videos for our face images. Due to different camera app implementation, we have to perform following transformation to get upright photos from video clips.

```
if className in ('XiaoYan', 'GabyNg'):
    image = cv2.transpose(image)
elif className in ('LeeSeng'):
    image = cv2.transpose(image)
    image = cv2.flip(image, 0)
```

5.2. The usage of jupyter notebook and code quality

Data project is still a bunch of source code. It is normal to write spaghetti code to try things out and show the result in jupyter notebook. However, it is critical for new developer to avoid using jupyter notebook at all cost, as it doesn't provide refactoring facility. Jupyter notebook is meant for generating research reports. And yet after trying parameters and network setup, all reported figures disappeared into 1 jupyter notebook, and replication of the result requires meddling with parameters again.

It would be better to use IDE with refactoring facility, e.g. pycharm community edition, to refactor spaghetti code into methods for reusability.



With simple methods, it will be easier to generate reports under jupyter notebook, and produce less errors in various reports under different scenario.

5.3. Slower training due to single process image augmentation

We always wonder why low GPU usage during slow training.

```

leeseng@DeepLearner: ~
File Edit View Search Terminal Tabs Help

leeseng@DeepLearner: ~
leeseng@DeepLearner: ~

Thu Sep 26 21:42:35 2019

=====
| NVIDIA-SMI 430.26             Driver Version: 430.26             CUDA Version: 10.2             |
|                               |                               |                               |
| GPU  Name                      Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|=====|=====|=====|
|  0  GeForce RTX 208...      Off   | 00000000:01:00:0 Off |             N/A             |
| 32%   50C   P2      60W / 250W | 7383MiB / 11016MiB |      8%      Default  |
|=====|=====|=====|
|  1  GeForce GTX 1080      Off   | 00000000:02:00:0 Off |             N/A             |
| 0%   40C   P8      6W / 200W  | 2MiB / 8119MiB   |     0%      Default  |
|=====|=====|=====|

Processes:
=====
| GPU  PID    Type   Process name                      GPU Memory |
|=====|=====|=====|
|  0    2075   G     /usr/lib/xorg/Xorg                30MiB     |
|  0    3757   G     /usr/bin/gnome-shell              15MiB     |
|  0    10924  C     /usr/bin/python3                  7325MiB   |
|=====|=====|=====|

```

activation_27 (Activation)	(None, 512)	0
dropout_16 (Dropout)	(None, 512)	0
dense_7 (Dense)	(None, 3)	1539
=====		
Total params: 6,355,011		
Trainable params: 6,355,011		
Non-trainable params: 0		

```
model summary: None
Found 4488 validated image filenames belonging to 3 classes.
Found 1306 validated image filenames belonging to 3 classes.
Learning rate: 0.001
Epoch 1/50
35/35 [=====] - 59s 2s/step - loss: 0.9921 - acc: 0.5498 - val_loss: 0.4096 - val_acc:
0.8508
Learning rate: 0.001
Epoch 2/50
35/35 [=====] - 58s 2s/step - loss: 0.2145 - acc: 0.9220 - val_loss: 0.0585 - val_acc:
0.9875
Learning rate: 0.001
Epoch 3/50
11/35 [=====>.....] - ETA: 27s - loss: 0.0405 - acc: 0.9865
```

```

leeseng@DeepLearner: ~
File Edit View Search Terminal Tabs Help

leeseng@DeepLearner: ~
leeseng@DeepLearner: ~

1 [| 0.7%] 5 [| 0.0%]
2 [| 0.0%] 6 [| 1.3%]
3 [| 0.7%] 7 [||||| 99.3%]
4 [| 0.0%] 8 [||||| 0.0%]
Mem[||||| 17.02G/31.4G] Tasks: 135, 486 thr; 2 running
Swp[ 0K/31.9G] Load average: 0.99 0.57 1.47
Uptime: 05:08:51

PID USER PRI NI VIRT RES SHR S CPU% MEM% TIME+ Command
10924 leeseng 20 0 20.9G 3643M 665M S 101. 11.3 3:17.24 /usr/bin/python3 -m
12817 leeseng 20 0 20.9G 3643M 665M R 100. 11.3 0:08.49 /usr/bin/python3 -m
12821 leeseng 20 0 27172 4440 3444 R 0.7 0.0 0:00.07 htop
9585 leeseng 20 0 8511M 1327M 9608M S 0.7 4.1 3:49.60 /home/leeseng/.local
11084 leeseng 20 0 17824 6564 3476 S 0.0 0.0 0:05.51 nvidia-smi -l
9645 leeseng 20 0 8511M 1327M 9608M S 0.0 4.1 0:03.06 /home/leeseng/.local
10994 leeseng 20 0 20.9G 3643M 665M S 0.0 11.3 0:01.19 /usr/bin/python3 -m
9126 leeseng 20 0 2294M 160M 8488M S 0.0 0.5 0:09.33 usr/share/jetbrains-
9598 leeseng 20 0 8511M 1327M 9608M S 0.0 4.1 0:02.04 /home/leeseng/.local
10999 leeseng 20 0 20.9G 3643M 665M S 0.0 11.3 0:02.24 /usr/bin/python3 -m
1 root 20 0 220M 9556 6684 S 0.0 0.0 0:01.79 /sbin/init splash
388 root 19 -1 147M 50156 49096 S 0.0 0.2 0:00.44 /lib/systemd/systemd
414 root 20 0 40024 6220 3072 S 0.0 0.1 0:01.24 /lib/systemd/systemd
F1Setup F2Search F4Filter F5Tree F6SortBy F7Nice F8Size F9All F10Quit

```

After observe CPU usage, we suspect ImageDataGenerator is single process by default.

After searching around tensorflow API, we noticed that Model.fit_generator() seems to be able to do multiprocessing.

- **workers** : Integer. Maximum number of processes to spin up when using process-based threading. If unspecified, **workers** will default to 1. If 0, will execute the generator on the main thread.
- **use_multiprocessing** : Boolean. If **True** , use process-based threading. If unspecified, **use_multiprocessing** will default to **False** . Note that because this implementation relies on multiprocessing, you should not pass non-picklable arguments to the generator as they can't be passed easily to children processes.

After setting “workers=5, use_multiprocessing=True”, we reduce the training time to 1/3 of existing training time, from 58seconds to 18 seconds. Suggests workers number to be “total cpu count - 2”. The 2 cpus are reserved for OS and jupyter notebook main process.

```
dropout_20 (Dropout)          (None, 512)          0
dense_9 (Dense)               (None, 3)            1539
=====
Total params: 6,355,011
Trainable params: 6,355,011
Non-trainable params: 0
model summary: None
Found 4524 validated image filenames belonging to 3 classes.
Found 1297 validated image filenames belonging to 3 classes.
Learning rate: 0.001
Epoch 1/50
35/35 [=====] - 18s 510ms/step - loss: 0.8433 - acc: 0.6015 - val_loss: 0.2785 - val_a
cc: 0.8852
Learning rate: 0.001
Epoch 2/50
35/35 [=====] - 18s 517ms/step - loss: 0.1446 - acc: 0.9570 - val_loss: 0.0770 - val_a
cc: 0.9859
Learning rate: 0.001
Epoch 3/50
35/35 [=====] - 18s 518ms/step - loss: 0.0377 - acc: 0.9879 - val_loss: 0.0203 - val_a
cc: 0.9969
Learning rate: 0.001
Epoch 4/50
8/35 [=====>.....] - ETA: 11s - loss: 0.0261 - acc: 0.9951
```

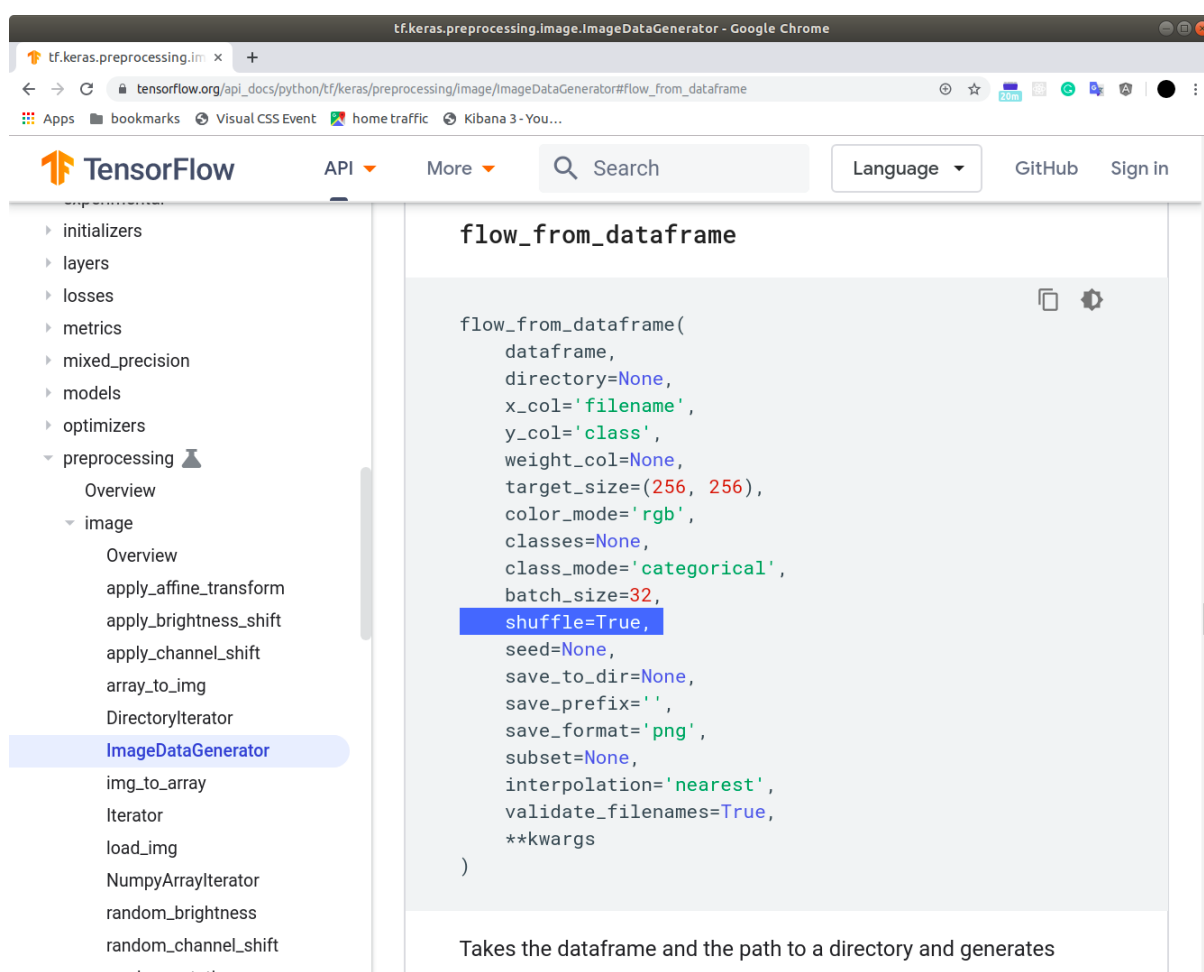
CPU and GPU usage also become more fully utilized.

```
leeseng@DeepLearner: ~  
File Edit View Search Terminal Tabs Help  
leeseng@DeepLearner: ~ x leeseng@DeepLearner: ~ x  
+-----+  
Thu Sep 26 21:47:56 2019  
+-----+  
| NVIDIA-SMI 430.26          Driver Version: 430.26          CUDA Version: 10.2          |  
+-----+  
| GPU Name Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |  
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |  
+-----+  
| 0  GeForce RTX 208...  Off | 00000000:01:00.0 Off |           N/A       |  
| 39%   65C   P2    97W / 250W | 7383MiB / 11016MiB |      40%     Default |  
+-----+  
| 1  GeForce GTX 1080    Off | 00000000:02:00.0 Off |           N/A       |  
| 0%    40C   P8     6W / 200W |   2MiB /  8119MiB |       0%     Default |  
+-----+  
+-----+  
| Processes:                                     GPU Memory |  
| GPU       PID    Type    Process name                               Usage       |  
+-----+  
|    0      2075     G   /usr/lib/xorg/Xorg                           30MiB       |  
|    0      3757     G   /usr/bin/gnome-shell                        15MiB       |  
|    0     10924     C   /usr/bin/python3                          7325MiB      |  
+-----+  
|
```

```
leeseng@DeepLearner: ~  
File Edit View Search Terminal Tabs Help  
leeseng@DeepLearner: ~ x leeseng@DeepLearner: ~ x  
1 [||| 4.0%] 5 [|||||94.7%]  
2 [|||| 13.2%] 6 [|||||94.7%]  
3 [|||||94.7%] 7 [||| 5.3%]  
4 [|||||94.1%] 8 [|||||94.7%]  
Mem[|||||9.32G/31.4G] Tasks: 155, 518 thr; 6 running  
Swp[||| 0K/31.9G] Load average: 1.09 0.32 0.60  
Uptime: 05:33:32  
  
PID USER PRI NI VIRT RES SHR S CPU% MEM% TIME+ Command  
13589 leeseng 20 0 22.6G 4645M 82148 R 93.9 14.5 0:01.42 /usr/bin/python3 -m  
13591 leeseng 20 0 22.6G 4644M 82148 R 93.2 14.5 0:01.41 /usr/bin/python3 -m  
10924 leeseng 20 0 22.6G 5207M 665M S 19.2 16.2 7:38.25 /usr/bin/python3 -m  
11000 leeseng 20 0 22.6G 5207M 665M S 9.9 16.2 0:05.05 /usr/bin/python3 -m  
13560 leeseng 20 0 22.6G 5207M 665M S 6.6 16.2 0:00.37 /usr/bin/python3 -m  
10995 leeseng 20 0 22.6G 5207M 665M S 2.0 16.2 0:07.72 /usr/bin/python3 -m  
13574 leeseng 20 0 27252 4684 3528 R 1.3 0.0 0:00.21 htop  
9585 leeseng 20 0 8511M 1327M 90680 S 0.7 4.1 4:06.89 /home/leeseng/.local  
9693 leeseng 20 0 8511M 1327M 90680 S 0.7 4.1 0:18.55 /home/leeseng/.local  
9710 leeseng 20 0 992M 65584 18556 S 0.7 0.2 0:09.29 /usr/bin/python3.6 /  
9598 leeseng 20 0 8511M 1327M 90680 S 0.7 4.1 0:03.25 /home/leeseng/.local  
10993 leeseng 20 0 22.6G 5207M 665M S 0.0 16.2 0:01.49 /usr/bin/python3 -m  
11084 leeseng 20 0 17824 6564 3476 S 0.0 0.0 0:09.98 nvidia-smi -l  
F1Help F2Setup F3Search F4Filter F5Tree F6SortBy F7Nice - F8Nice + F9Kill F10Quit
```


5.4. ImageDataGenerator.flow_from_dataframe(), implicit shuffle=True

This is one of the tricky bugs to catch under testing. The true labels are provided by dataframe, but `flow_from_dataframe()` API has “`shuffle = True`” by default. We can only observe random accuracy even with same model for a few rounds of testing. We discover that by applying the validation set images to testing. With validation accuracy of 90% during training, it is weird that we got random accuracy testing score again. After recalling that “`shuffle = True`” in `ImageDataGenerator.flow_from_dataframe()`, setting “`shuffle=False`” resolve this testing bug.



5.5. High validation accuracy BUT Low testing accuracy

During our initial training both our model using full images for both training and validation, the validation accuracy is 100%. But when we apply the model for test images prediction, the accuracy is very low, most of the models get accuracy less than 60%. The likely possible reason is due to larger and similar background of faces during selfie video. Therefore, we learn the hard way only during testing phase. Fortunately, we have [face_recognition](#) API to crop faces for training and testing, which helps to improve our final result accuracy to 85%, from 67%.

6. Conclusion

Using a CNN model, we started our training and testing using the 2D images we extracted from our videos. After training the model with different parameters and we could not improve the accuracy above 50%, we realised the 2D images extracted from our video, contains too much background “noise” to the face images. As such, we determined that in order to get a better accuracy of our face recognition, it is important to extract only the face images to train our model.

We got the worst test accuracy for full images especially wrong prediction for the child’s images. The child’s full images’ background maybe is the noise for the training which cause the prediction not accurate problem. That prompts us to consider cropping only face for training. We researched on the Internet and found https://github.com/ageitgey/face_recognition API about face clipping. To remove the noise of the images, we use face crop images for our models training, the test accuracy much better and achieve 85%.

References:

<https://medium.com/@ageitgey/machine-learning-is-fun-part-4-modern-face-recognition-with-deep-learning-c3cffc121d78#.ds8i8oic9>